Reducing the Barrier to Entry of Complex Robotic Software: a MoveIt! Case Study

David COLEMAN[1]   Ioan ŞUCAN[2]   Sachin CHITTA[3]   Nikolaus CORRELL[1]

[1] Dept. of Computer Science, University of Colorado at Boulder, 430 UCB, Boulder, CO 80309
[2] Google Inc., 1600 Amphitheatre Parkway, Mountain View, CA 94043
[3] SRI International, 333 Ravenswood Avenue, Menlo Park, CA 94025

Abstract—Developing robot agnostic software frameworks involves synthesizing the disparate fields of robotic theory and software engineering while simultaneously accounting for a large variability in hardware designs and control paradigms. As the capabilities of robotic software frameworks increase, the setup difficulty and learning curve for new users also increase. If the entry barriers for configuring and using the software on robots is too high, even the most powerful of frameworks are useless. A growing need exists in robotic software engineering to aid users in getting started with, and customizing, the software framework as necessary for particular robotic applications. In this paper a case study is presented for the best practices found for lowering the barrier of entry in the MoveIt! framework, an open-source tool for mobile manipulation in ROS, that allows users to 1) quickly get basic motion planning functionality with minimal initial setup, 2) automate its configuration and optimization, and 3) easily customize its components. A graphical interface that assists the user in configuring MoveIt! is the cornerstone of our approach, coupled with the use of an existing standardized robot model for input, automatically generated robot-specific configuration files, and a plugin-based architecture for extensibility. These best practices are summarized into a set of *barrier to entry design principles* applicable to other robotic software. The approaches for lowering the entry barrier are evaluated by usage statistics, a user survey, and compared against our design objectives for their effectiveness to users.

Index Terms—Robotic Software Frameworks, Motion Planning, Barrier to Entry, Setup, Usability, MoveIt!

1 INTRODUCTION

MANAGING the increasing complexity of modern robotic software is a difficult engineering challenge faced by roboticists today. The size of the code bases of common open source robotic software frameworks such as ROS [1], MoveIt! [2] and OROCOS [3] continues to increase [4], and the required breadth of knowledge for understanding the deep stack of software from control drivers to high level planners is becoming more daunting. As it is often beyond the capabilities of any one user to have the necessary domain knowledge for every aspect of a robot's tool chain, it is becoming increasingly necessary to assist users in the configuration, customization, and optimization of the various software components of a reusable robotic framework.

The user interface design principles required in the emerging field of robotics software is similar to other more mature software engineering fields and much can be learned from them. There have been many examples of software, such as computer operating systems, that have historically required many installation and configuration steps whose setup process has since improved. Still, the user interface design principles for robotics is unique in 1) the degree to which software interacts with hardware and real world environments compared to consumer-level software, 2) the large variety in complexity and scale of robotic platforms, and 3) the long term desire to increase the autonomy of robotics systems by reducing reliance on GUIs and increasing high level robotic intelligence.

1.1 Barriers to Entry

The term *barriers to entry* is used in the context of robotic software engineering to refer to the time, effort, and knowledge that a new user must invest in the integration of a software component with an arbitrary robot. This can include, for example, creating a virtual model of the robot's geometry and dynamics, customizing configuration files, choosing the fastest

algorithmic approach for a specific application, and finding the best parameters for various algorithms.

Powerful robotics software generally requires many varying degrees of customization and optimization for any particular robot to operate properly. Choosing the right parameters for each utilized algorithm, software component, and application typically involves expert human input using domain-specific knowledge. Many new users to a software package, particularly as robotics becomes more mainstream, will not have the breadth of knowledge to customize every aspect of the tool chain. When the knowledge of a new user is insufficient for the requirements of the software, the barriers to entry become insurmountable and the software unusable. One of the emerging requirements of robot agnostic frameworks is implementing mechanisms that will automatically setup and tune task pipelines for arbitrary robots.

Another motivation for lowering the barrier to entry of complex robotics software is the *paradox of the active user*. This paradox explains a common observation in many user studies that *users never read manuals* but start attempting to use the software immediately [5]. The users's desire to quickly accomplish a task results in their skipping the reading of any provided documentation or gaining deeper understanding of the system and instead diving right into completing their task. The *paradox* is that the users would actually save time in the long run if they learned more about the system before attempting to use it, but these studies showed that in reality people do not tend to invest time upfront into learning a new system.

Even experts in the area of the associated robotics software will become frustrated with robotics software if all initial attempts to setup and configure the framework fail and no progress is made. Most researchers and engineers typically do not have the time or ability to completely understand the entirety of robotics software before they start using it. It is important for the user's initial experience with a piece of software to be positive to ensure its continued use.

1.2 Benefits of Larger User Base

The need to lower the barrier of entry is beneficial to the software itself in that it enables more users to utilize the framework. If the software framework is being sold for profit, the benefits of a larger user base are obvious. If instead the software is a free open-source project, as many successful robotic frameworks currently are [4], lowering the barrier to entry is very beneficial in that it creates the *critical mass of skilled contributors* that has been shown to make open source projects successful [3]. As the number of users increases, the speed in which bugs are identified and fixed increases [6]. It is also typically hoped that development contributions to the code base increases, though this correlation is not as strong [6]. One of the key strengths of a larger community for an open source project is increased participation of users assisting with quality assurance, documentation, and support [7].

Another benefit of lowering the barrier of entry is that it allows the robotics software to become an educational tool for robotics. Not only is the software accessible for academic research and industrial applications, but graduate, undergraduate, and even primary-level students can use it to learn some of the higher level concepts of robotic applications as has been demonstrated in [8], [9], [10].

Beyond the motivation of success for an individual software project, broadening access to robotics software development increases the number of creative minds working on solving today's challenging robotics problems. Making the accessibility of robotic development more like mobile device development and web development might increase the speed of innovation in robotics, similar to that experienced by phone apps and the Internet [11].

As used in this paper, the target *users* are engineers, scientists, students, and hobbyists with a general aptitude for software and robotics, but who are not necessarily experts in either of those fields. The hope remains that human-robotic interaction for the general population in the future will be based on more natural methods and that software configuration and graphical user interfaces (GUIs) are only necessary for the robot developers themselves [12].

1.3 Difficulty of Robot Agnostic Software

The software engineering challenges faced in making reusable, robot agnostic code are hard and are different than those in other re-usable software frameworks. In [13], Smart argues there are three main factors that make general-purpose robotic software difficult: heterogeneity of robotics, limited resources (computational or otherwise), and the high rate of hardware and software failures. The variety of different tasks and task constraints imposed upon robots is another challenge for robot agnostic software [14].

The heterogeneity of robots is of primary concern to us in this paper – accounting for different types of actuators, sensors, and overall form factors is a very difficult task. To some users a robot is a robust and precise industrial arm, to others a robot is simply a mobile base with wheels and a computer, and to others a robot is a fully anthropomorphic biped. Creating reasonable abstractions for these large amounts of variation requires many trade offs to be made that almost always lead to a sub-optimal solution for all robots. It is more difficult to create a hardware abstraction for a robot than a standard computer – when robotic software must interact with physical devices through an abstraction, it gives up specific knowledge of the hardware that then requires much greater reasoning and understanding of its configuration [13].

Operational requirements are another challenge in making reusable software for a wide range of robotic platforms. Some users require hard real-time constraints, while others can tolerate "fast enough" or "best effort" levels of performance. Variable amounts of available computational resources such

as processing power or the "embeddedness" of the system also makes it difficult to design robot agnostic code that can run sufficiently on all robots. The amount of required error checking and fault tolerance varies by application area, for example, there are significant differences between a university research robot and a space exploration rover or a surgical robot.

In making robotic agnostic software, many time-saving shortcuts employed for single-robot software must be avoided. This includes hard coding domain-specific values for "tweaking" performance and using short-cutting heuristics applicable to only one hardware configuration. Instead reasonable default values or automatically optimized parameters must be provided as discussed later.

On top of these challenges, packaging reusable software into an easy to setup experience for end users requires creating tools that automate the configuration of the software.

1.4   Related Work

There has been much work to address the software engineering challenges of complex robotic frameworks, but typically the identified design goals have emphasized the need for features such as platform independence, scalability, real-time performance, software reuse, and distributed layouts [15], [16], [17]. In [17]'s survey of nine open source robotic development environments, a collection of metrics was used which included documentation and GUIs, but no mention was made of setup time, barrier to entry, or automated configuration.

A focus on component-based design of motion planning libraries similar to MoveIt! was addressed in [18]. The challenges of software reuse, combining various algorithms, and customizations are discussed, but the work falls short of addressing the initial ease of use of these robotics frameworks. The difficulty of creating good component abstractions between hardware and algorithms is addressed in [14].

The importance of an open source robotics framework having a large number of researchers and engineers motivated to contribute code and documentation is emphasized in the OROCOS framework [3], which we emphatically agree with, but take a step further by creating additional tools to encourage higher user adoption.

Human-robot interaction (HRI) has also been a popular area of research, but HRI's focus has been on the runtime behavior of robots and not on the difficulties of human users applying software frameworks to robot hardware [19], [12], [20]. For example, in [21], an effective user interface is presented for teleoperation of rescue robots, but no thought is given to making it robot agnostic or to its configuration.

In [22], Chitta et. al. we presented a set of tools that allowed the Arm Navigation software framework (the precursor to MoveIt!) to be easily configured within a short amount of time for a new robotic system. This paper extends and improves that work, focusing specifically on the difficulties of setting up and configuring robotics software.

1.5   Contribution and Outline

In this paper, we will present best practice principles for lowering the barrier of entry to robotic software using the new MoveIt! software [2] as our case study. In section 2 we will motivate the many software components that make software like MoveIt! a good example of the difficulties of complex robotics software. In section 2.2, we will briefly describe MoveIt! itself. In section 3 we explain the design principles used to address the user interface needs of robotics software. We then in section 4 show how we have taken these principles and implemented them in an entry tool for MoveIt! to reduce the entry barriers. In section 5, we will present the results of these implementations on the size of the user base and ease of adoption of the MoveIt! software framework. We will discuss our experiences and lessons learned in section 6, followed by our conclusion in section 7.

2   MOTION PLANNING FRAMEWORKS

Robotic motion planning is a maturing and central field in robotics [9] that turns a high level task command into a series of discrete motions such that a robot can move within its environment. The typical use case considered in this paper is the problem of controlling a robotic arm from one configuration to another while taking into account various constraints.

The software development of a motion planning framework (MPF) is challenging and involves combining many disparate fields of robotics and software engineering [23]. We refer to the software as a *framework* in this context because it abstracts the various components of motion planning into generic interfaces as discussed later.

One of the most important features of a MPF is providing the structures and classes to share common data between the different components. These basic data structures include a model of the robot, a method for maintaining a representation of the state of the robot during planning and execution, and a method for maintaining the environment as perceived by the robot's sensors (the "planning scene").

In addition to the common data structures, a MPF requires many different interacting software components, henceforth referred to as the *planning components*. A high level diagram of the various planning components is shown in Figure 1. The planning component that actually performs motion planning includes one or more algorithms suited for solving the expected problems a robot will encounter. The field of motion planning is large, and no one-size-fits-all solution exists yet, so a framework that is robot agnostic should likely include an assortment of algorithms and algorithm variants.

Other planning components include a collision checking module that detects the potential intersection of geometric primitives and meshes in the planning scene and robot model. A forward kinematics solver is required to propagate the robot's geometry based on its joint positions, and an inverse kinematics solver is required when planning in the Cartesian
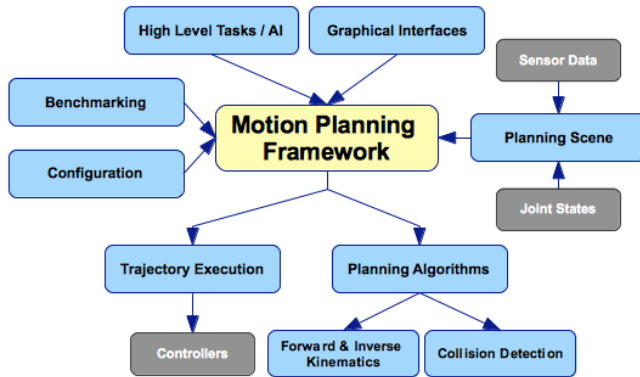
Fig. 1. High level diagram of various planning components (blue boxes) in a Motion Planning Framework (MPF). Grey boxes represent external input and output.

space of the end effector for some planning techniques. Other potential constraints, such as joint/velocity/torque limits and stability requirements, require additional components.

Secondary components must also be integrated into a powerful MPF. Depending on what configuration space a problem is solved in, the generated motion planning solution of position waypoints must be parameterized into a time-variant trajectory to be executed. A controller manager must decide the proper low level controllers for the necessary joints for each trajectory. A perception interface updates the planning scene with recognized objects from a perception pipeline as well as optional raw sensor data.

Higher level applications are built on top of these motion planning components to coordinate more complex tasks, such as pick and place routines. Other optional components of a MPF can include benchmarking tools, introspection and debugging tools, as well as the user-facing GUI.

2.1 Existing Motion Planning Software

Many open source software projects for motion planning exists whose intent is to provide a platform for testing and developing novel path planning algorithms and other motion planning components. We will distinguish them from a motion planning *framework* due to their exclusion of actual hardware perception and control. All offer varying degrees of modularity and all have a basic visualization window for viewing motion plans of 3D geometries. A brief review of them is presented here.

Both LaValle's Motion Strategy Library (MSL) [24], 2000, and Latombe's Motion Planning Kit (MPK) [25], 2003, have scopes limited to only simulation and therefore are not frameworks in our definition. The MSL is configured manually using six required text files and up to fifteen optional files, depending on the planning problem. It has a GUI for tweaking parameters and controlling the visualization of plans. The MPK is able to load robots with varying geometry without recompiling code and provides a scene format that is an extension of the Open

Inventor format. It does not have a fully interactive GUI but rather allows control only through keyboard shortcuts. Neither MSL or MPK provides assistance for setting up a new robot and has little to no documentation on this process.

The Karvaki Lab's Object-Oriented Programming System for Motion Planning (OOPSMP) [26], 2008, is a predecessor to the Open Motion Planning Library (OMPL) [27], 2010, both of which are collections of planning algorithms and components whose scope also excludes hardware execution and perception tasks. OOPSMP is XML based for configuration, scene definitions, and robot geometry. An additional SketchUp interface provides a quick way to build environments. It has some GUIs that assist in visualization. OMPL differs in its handling of environments and robots in that it abstracts that notion into a black box and instead operates in various configuration spaces. Neither OOPSMP nor OMPL provides any tools or GUIs for configuration.

Diankov's OpenRave [28], 2010, is a fully featured motion planning framework with many high level capabilities, some GUIs, and the ability to connect to hardware controllers and sensors. It uses the Collada format [29], as well as its own proprietary format, to define robots and environments. Its main interface is through simple python scripting, and it utilizes a plugin interface to provide extensibility of the framework. It too falls short of providing easy to setup tools for new robots.

Willow Garage's ROS Arm Navigation framework [22], 2010, is the predecessor of MoveIt! and provides much of the same functionality of MoveIt! and OpenRave but also includes a Setup Wizard that provides a GUI for helping new users setup arbitrary robots into the framework. It was the inspiration for the Setup Assistant described later in this paper.

2.2 MoveIt! Motion Planning Framework

MoveIt![2] is the primary software framework for motion planning and mobile manipulation in ROS and has been successfully integrated with many robots including the PR2 [30], Robonaut [31], and DARPA's Atlas robot. MoveIt! is written entirely in C++ but also includes Python bindings for higher level scripting. It follows the principle of software reuse as advocated for robotics in [4] of not tying itself exclusively to one robotic framework—in its case ROS—by creating a formal separation between core functionality and robotic framework-dependent aspects (e.g., communication between components).

MoveIt! uses by default the core ROS build and messaging systems. To be able to easily swap components MoveIt! uses plugins for most of its functionality: motion planning plugins (currently using OMPL), collision detection (currently using the Fast Collision Library (FCL) [32]), kinematics plugins (currently using the OROCOS Kinematics and Dynamics Library (KDL) [33] for forward and inverse kinematics for generic arms as well as custom plugins). The ability to change these default planning components is discussed in section 4.3. MoveIt!'s target application is manipulation (and

mobile manipulation) in industrial, commercial and research environments. For a more detailed description of MoveIt!, the interested reader is referred to [2].

3 ENTRY BARRIER DESIGN PRINCIPLES

In designing the configuration process that enables MoveIt! to work with many different types of robots, with almost any combination of planning components, several contending design principles for lowering the barrier of entry emerged. These requirements were drawn partially from standard HCI principles [34], from work on MoveIt!'s predecessor, and from an iterative design process where feedback was gained from internal users at Willow Garage during development. We believe these *entry barrier design principles* transcend motion planning and can be applied to most robotic software:

Immediate: The amount of time required to accomplish the most primitive task expected from the robotic software component should be minimized. This is similar to the time-honored "Hello World" demo frequently used by programming languages and typical Quick Start guides in documentation. Immediacy is essential for the *paradox of the active user* as it provides cursory feedback to the user that the software works and is worth investing further time.

Transparent: The configuration steps being performed automatically for the user, and the underlying mechanisms utilized in the software components, should be as visible as possible. Transparency is important so that users can later understand what parameters are specific to their robot and know how to customize the aspects they desire. A "layered" approach of presenting information can offer a good balance of separating the required knowledge for a user's immediate goals from the "useful later" information needed to prevent the user from being hindered in the future.

Intuitive: The need to read accompanied documentation, and the amount of required documentation, should be minimized. A well-designed user interface, be it graphical or command line, should be as intuitive as possible by following standard design patterns and providing interface context clues. An ideal GUI for configuration would not require any documentation for most users.

Reconfigurable: The automatically generated parameters and default values for the initial setup of a robot should be easy for the user to modify at a later time. Typically, these parameters and values are chosen to work for the largest number of robots possible but are not optimal for any particular robot. Providing easy methods to reconfigure the initial setup is important for allowing better performance.

Extensible: The user should be enabled to customize as many components and behaviors as possible within the reasonable scope of the software. Providing the means to extend the software with custom solutions for a particular application makes the software far more powerful and re-usable for varying use-cases. A typical solution for this is providing a plugin interface.

Documented: The amount of reference material explaining how to use the software should be maximized for as many aspects and user levels as possible. Even the most intuitive software requires documentation for various aspects of the operation or modification of the software itself. Different types of documentation are needed for different users—for example developers and end users—though in Robotics these groups are frequently the same. Documentation is arguably the most important factor in reducing the barrier to entry of new software [35].

These principles are additionally applicable to computer software in general, but a greater focus on hardware variance and the needs of developers has been applied. *Reconfigurability*, or personalization, is common in computer software as well, but in our application we use it mainly in reference to parameters that require customization for different physical geometries and hardware designs. Similarly, *extensibility* and *transparentness* are design principles that aid robotic developers in applying their software to specific hardware. Whereas today most computer hardware is fairly standardized and share similar capabilities, robotics hardware still has large variability in design and capability. For this reason, *transparency* is particularly important since many robotic researchers and developers need to understand the software enough to adapt it to their unique hardware.

Many of these *entry barrier design principles* have opposing objectives that require a balance to be found between them. For example, the desire for *transparency* in the underlying mechanisms often leads to slower setup times (lack of *immediacy*) and more complicated configuration steps (lack of *intuitiveness*). The need for extensibility of various components in the software often results in far more complicated software design as more abstraction is required, resulting in a less *intuitive* code base and difficult *documentation*. Nevertheless, compromises can be made between these principles that result in a superior user experience, as will be demonstrated in the next section.

4 METHODS TO LOWER THE ENTRY BARRIER

One of the unique features of MoveIt! is the ratio of its power and features to the required setup time. A beginner to motion planning can take a model of their robot and with very little effort execute motion plans in a virtual environment. With a few additional steps of setting up the correct hardware interfaces, one can then execute the motion plans on actual robotic hardware.

The *entry barrier design principles* discussed above were applied to MoveIt! to address the challenges faced for new users to this complex software framework. Developing these solutions required difficult software challenges to be overcome as discussed in the following case study.

4.1   Basic Motion Planning Out of the Box

To address the entry barrier design principle of *immediacy*, a streamlined "Quick Start" for MoveIt! was created that consists of a series of fairly trivial steps, relative to our target users. The most challenging of these steps—creating a *robot model*—is not directly related to the configuration of MoveIt! but rather is a prerequisite of using the software framework. Nevertheless, we will discuss this important prerequisite before proceeding to the more directly-related configuration steps.

Robot Model Format: The robot model is the data structures and accompanying file format used to describe the three-dimensional geometric representation of a robot, its kinematics, as well as other properties relevant to robotics. These other properties can include the geometric visualization meshes, courser-grained collision geometry of the robot used for fast collision checking, joint limits, sensors, and dynamic properties such as mass, moments of inertia, and velocity limits. Often the robot's joints and links relationships are represented by a kinematic tree, though this approach is problematic when a robot has a closed chain. In our application, as well as most state of the art MPFs, we will restrict our definition of modeled robots to arbitrarily articulated rigid bodies.

Extensible robotics software requires using a standardized format that can express the intricacies of varying hardware configurations. An additional design requirement for this standardized format is that it is *intuitive* for users to setup. There are a few options for representing robots, and in MoveIt! it was accomplished by using the Unified Robotic Description Format (URDF [36]) Document Object Model. This data structure is populated by reading human-readable (*transparent*) XML schemas – both URDF-formatted files (different from the datastructure) as well as the industry standard Collada [29] format.

Creating an accurate model of a robot can be a difficult task. URDF models for many robots already exist, so often users can avoid this problem. However, when a custom robot requires a new robot model, the URDF model in ROS was found to be the most appropriate to use since the user can also take advantage of tools in ROS for working with the URDF. In particular, there are tools for verifying the validity of the XML, for visualizing it, and for converting a SolidWorks CAD model of a robot directly into URDF format.

MoveIt! Setup Assistant: The main facility that provides out of the box support for beginners is the MoveIt! Setup Assistant (SA). The SA is a GUI that steps new users though the initial configuration requirements of using a custom robot with the motion planning framework (Figure 2). It accomplishes the objective of *immediacy* for the user by automatically generating the many configuration files necessary for the initial operation of MoveIt!. These configurations include a self-collision matrix, planning group definitions, robot poses, end effector semantics, virtual joints list, and passive joints list.
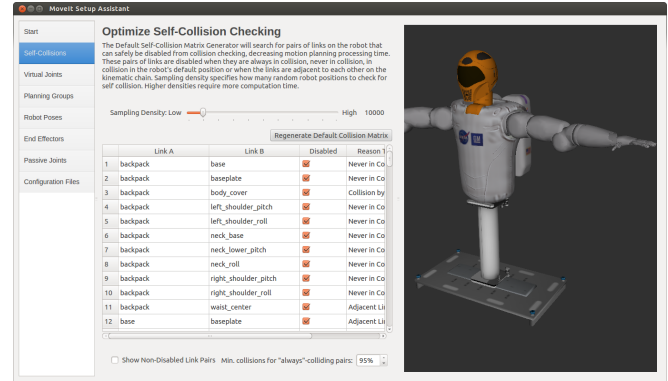

Fig. 2.   MoveIt! Setup Assistant GUI with the NASA Robonaut loaded on the self-collision matrix screen.
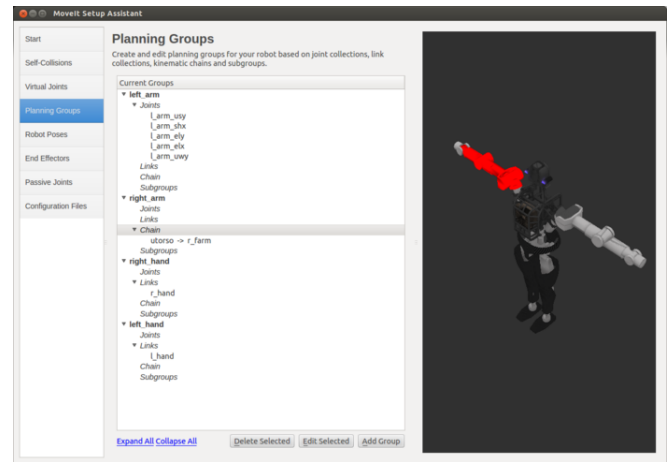

Fig. 3. MoveIt! Setup Assistant GUI with the Atlas robot's left arm highlighted for user feedback on the planning groups screen.

The GUI consists of 1) a large navigation pane on the left that allows the user to move back and forth through the setup process as needed (providing quick *reconfigurability*), 2) the middle settings window that changes based on the current setup step being performed by the user, and 3) a right side visualization of the three dimensional model of the robot as it is being configured. The right side visualization increases the *immediacy* of results and *transparency* of the configuration by highlighting various links of the robot during configuration to visually confirm the actions of the user, as shown in Figure 3.

Using a properly formatted robot model file with the SA, MoveIt! can automatically accomplish many of the required tasks in a MPF. If one desired, the steps within the SA could almost entirely be automated themselves, but they have been kept manual so to 1) increase *transparency* and 2) provide *extensibility* for edge cases and unusual customizations. For example, automated semantical guesses of where an arm ends and an end effector begins can sometimes be incorrect.
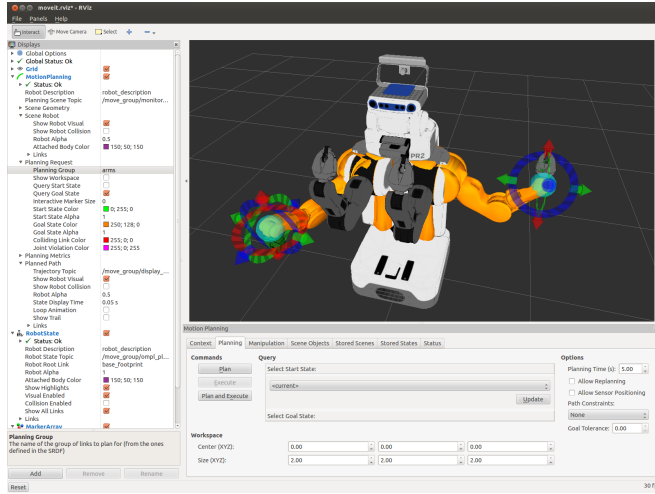
Fig. 4. MoveIt! Motion Planning Visualization GUI with the PR2 planning with both arms to goal positions with interactive mouse-based tools

Fig. 5. MoveIt! Motion Planning Visualization GUI with the Baxter robot visualizing steps of a motion plan

MoveIt! Motion Planning Visualization GUI: The details of the automated configuration are left for the next section, but after the steps in the SA are completed a demo script is created that automatically starts up a visualization tool with the new robot loaded and ready to run motion planning algorithms in a non-physics based simulation. A typical demo task would be using the computer mouse to visually drag 3D interactive arrows situated on the robot's end effector from a start position to a goal position around some virtual obstacle. The demo can then quickly plan the arm in a collision free path around the obstacle and visualize the results within the GUI.

This user interaction is accomplished with the MoveIt! Motion Planning Visualization (MMPV) [2], an additional GUI that allows beginning users to learn and experiment with a large subset of the functionality provided by MoveIt! (Figure 4). While the long term goal of robotics is to provide more autonomous solutions to motion planning and human-robot interactions [12], the MMPV fulfills the immediate needs of direct operation for testing and debugging the framework's capabilities easily. This interface is a vital component of MoveIt!'s strategy to provide *immediate* results for motion planning with a robot that does not require any custom coding. Once the user is comfortable with the basic feature set and functionality of MoveIt!, *extensibility* is provided via varying levels of code APIs for more direct, non-GUI, access to the robot's abilities.

The MMPV provides a large number of features and visual tools for motion planning. Using the MMPV, visualizations such as seen in Figure 5 are provided of:

- Start and goal configurations of the robot for planning
- Current robot hardware configuration
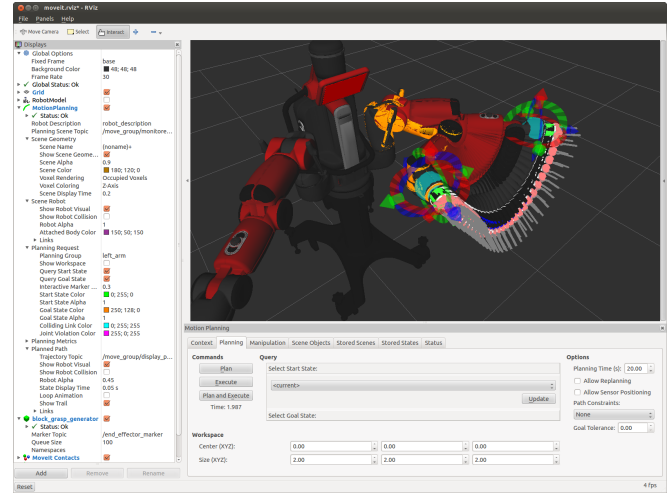- Animated planned path before execution
- Detected collisions

- Sensor data and recognized objects
- Pick and place data such as grasp positions
- Attached bodies such as manipulated objects
- Planning metrics

Additionally, the MMPV contains many other non-visualization tools such as:

- Connecting to a database of planning scenes
- Adjusting inverse kinematic settings
- Changing the utilized planning algorithm
- Adjusting the workspace size
- Adjusting goal tolerance and planning time
- Tweaking manipulation plans
- Loading and moving collision objects
- Exporting/importing scenes and states
- Viewing the status of MoveIt!

Hardware Configuration and Execution: Once the user is comfortable with the basic tools and features provided by MoveIt!, the next step is to configure their robot's actual hardware actuators and control interfaces to accept trajectory commands from MoveIt!. This step is not as easy and requires some custom coding to account for the specifics of the robot hardware—the communication bus, real-time requirements, and controller implementations. At the abstract level, all MoveIt! requires is that the robot hardware exposes its joint positions and accepts a standard ROS trajectory message containing a discretized set of time-variant waypoints including desired positions, velocities, and accelerations.

4.2 Automatic Configuration and Optimization

The size and complexity of a feature-rich MPF like MoveIt! requires many parameters and configurations of the software be automatically setup and tuned to improve the MPF's *immediacy*. MoveIt! accomplishes this in the 1) setup phase of

a new robot, using the Setup Assistant, 2) during the runtime of the application, and 3) using benchmarking and parameter sweeping[37].

Self-Collision Matrix: The first step of the SA is the generation of a self-collision matrix for the robot that is used in all future planning to speed up collision checking. This collision matrix encodes pairs of links on a robot that never need to be checked for self-collision due to the kinematic infeasibility of there actually being a collision. Reasons for disabled collision checking between two links includes:

- Links that can never intersect due to the reachability kinematics of the robot
- Adjacent links that are connected and so are by design in collision
- Links that are always in collision for any other reason, including inaccuracies in the robot model and precision errors

This self-collision matrix is generated by running the robot through tens of thousands of random joint configurations and recording statistics of each link pair's collision frequency. The algorithm then creates a list of link pairs that have been determined to never need to be collision checked. This reduces future motion planning runtimes because it reduces the amount of required collision checks for every motion planning problem. The algorithm is incomplete because in probabilistically rare cases a pair of links will be disabled for collision checking when they should not be. For this reason, the number of tests needs to be very high.

Configuration Files: The other six steps of the SA all provide graphical front ends for the data required to populate the Semantic Robotic Description Format (SRDF) and other configuration files used by MoveIt!. The SRDF provides *reconfigurable* semantic meta data of the robot model. It is data useful to motion planning but not relevant to the URDF because it does not describe physical properties of the robot. The SRDF information includes which set of joints constitutes an arm and which set of links is considered part of the end effector. It is one of the main components that allows MoveIt! to be robot agnostic and to avoid dependencies on specific robots [2]. Requiring the user to configure all the semantic information by hand in a text editor would be tedious and more difficult than using a GUI. The GUI populates the available options for each input field in list boxes and guides the user through filling in the necessary fields with buttons and graphical feedback.

The last step of the SA is to generate all launch scripts and configuration files. This step outputs to file the information collected from the user during the step-by-step user interface, as well as generates a series of default configuration and launch scripts that are automatically customized for the particular robot using the URDF and SRDF information. These defaults include velocity and acceleration limits for each joint, kinematic solvers for each planning group, available planning algorithms, and projection evaluators for planning. Default

planning adapters are setup for pre- and post-processing of motion plans. Default benchmarking setups, controller and sensor manager scripts, and empty object databases are all generated using launch scripts, which essentially allow one to start different sets of MoveIt! functionality that are already put together.

These configuration files can easily be modified later from their default values by simply editing the text-based configuration files. The format of the files are based on ROS standards, which were chosen for their wide spread acceptance, readability, and simplicity. For the launch files an XML-based format custom to launching ROS applications was utilized. For all other configuration files the open source YAML data serialization format was used.

Automatic Runtime Tuning: MoveIt! is designed to simplify solving planning problems by reducing the number of hard-coded parameters and so called "magic numbers." Sampling-based planning algorithms in particular require such parameters as input. MoveIt! uses heuristics from OMPL to automatically choose good values for certain parameters to reduce the amount of expert domain knowledge required and make MoveIt! extensible to a larger set of problems.

An example of automatic runtime tuning is the resolution at which collision checking is performed; it is defined as a fraction of the space extent. The space extent is the lowest upper bound for the distance between the two farthest configurations of the robot. This distance depends on the joint limits, the types of joints, and the planning groups being used. Using the same information, projections to Euclidean spaces can also be defined. These projections are used to estimate coverage during planning. For example, the projections for robot arms are orthogonal ones, using the joints closer to the shoulder, as those most influence the position of the end-effector.

Benchmarking: For applications that require more tuning and optimization than those afforded by automatically generated parameters and default values, MoveIt! provides the ability to configure and switch out different planning components and and specify their configuration. However, this capability is much less useful without the ability to quantify the results of different approaches. Optimization criteria such as path length, planning time, smoothness, distance to nearest obstacle, and energy minimization need benchmarking tools to enable users and developers to find the best set of parameters and planning components for any given robotic application.

MoveIt! lowers the barrier to entry to benchmarking by providing a command line-based infrastructure and benchmarking configuration files that allows each benchmark to easily be set up for comparison against other algorithms and parameters [37]. An additional GUI is currently in development that makes benchmarking easier, more *intuitive*, and reduces the learning curve to this feature set of MoveIt!.

Choosing the best combination of planning components and parameters for any particular robot and problem is a daunting task even for experts because of the number of choices
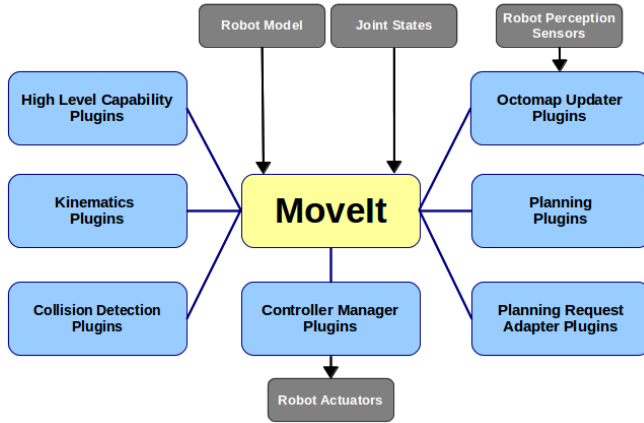
Fig. 6. Available planning component plugins for easily extending the functionality of MoveIt!. Grey boxes represent external input and output.

that must be made[37]. A common method to optimize an algorithms performance is to perform single and multivariable parameter sweeps during benchmarking. MoveIt! provides an interface for this in its benchmarking infrastructure by allowing an upper, lower, and increment search values to be provided by the user. Results can be output into generic formats for use in different plotting tools for analysis of which combination of parameters performed the best.

Attempting to fine-tune the functionality of MoveIt! with benchmarking and parameter sweeping is a feature for expert users and it is generally not required for entry-level users.

4.3 Easily Customize Framework Components

MoveIt! lowers the barrier to entry by not requiring users to provide their own implementation of any of the components in the motion planning framework. The default planning components are based on OMPL, FCL, and KDL. However, these default components are limiting to more advanced users who have their own application or research-specific needs to fulfill. We will briefly describe here how MoveIt! uses a plugin-based architecture and a high-level interface to address these *extensibility* issues (interested readers should refer to [2] for more detailed explanations).

Plugins: MoveIt! is designed to be *extensible* by allowing its various planning components to be customized through a lightweight plugin interface [2]. This is accomplished by using C++ shared objects that are loaded at run time, reducing dependency complexities. This plugin-centric framework, as seen in Figure 6, provides interfaces for forward and inverse kinematics, collision detection, planning, planning request adapters, controllers, perception, and higher level capabilities. Almost all aspects of MoveIt!'s functionality can be extended using plugins.

A particular strongpoint of MoveIt!'s feature set is its kinematics plugins that it can automatically generate using

the input URDF. The default KDL plugin uses numerical techniques to convert from a Cartesian space to joint configuration space. A faster solution can be achieved for some robots by utilizing OpenRave's IKFast [38] plugin that analytically solves the inverse kinematics problem. A combination of MoveIt! scripts and the IKFast Robot Kinematics Compiler can automatically generate the C++ code and plugin needed to increase the speed of motion planning solutions by up to three orders of magnitude [38].

Essentially, MoveIt! provides a set of data sharing and synchronization tools, sharing between all planning components the robot's model and state. The *extensibility* of MoveIt!'s framework is greatly enhanced by not forcing users to use any particular algorithmic approach.

High Level Interfaces: High level custom task scripting is easily accomplished in MoveIt! with both a C++ and python interface that abstracts away most of the underlying mechanisms in MoveIt!. Users who do not wish to concern themselves with how the various low level planning components are operating can focus instead on the high level application tasks, such as picking up an object and manipulating it. Python in particular is a very easy scripting language that enables powerful motion planning tasks to be accomplished with very little effort.

4.4 Documentation and Community

Though common place in open source software projects [3], it should be mentioned for completeness that MoveIt! addressed the *entry barrier design principle* of *documentation* by providing extensive online wiki pages, a mailing list for questions, and a issue tracker for bug reports and feature requests.

5 RESULTS

The success of MoveIt!'s efforts to lower its barrier of entry to new users through the application of the barrier to entry principles is quantified in the following. Its adoption rate, community activity, contributors, and results from a user survey are used as indicators of its progress.

5.1 Statistics

MoveIt! was officially alpha released on May 6th, 2013 — about 11 months prior to this writing. One method to quantify its popularity is by the total number of binary and source code installations that have been performed. Though not exactly representative of this data, MoveIt!'s website has an "Installation" page that receives an average of 940 unique page views per month [2] – a large fraction of that number can be assumed to represent unique installations.

There are currently 312 members on the MoveIt! mailing list as shown over time in Figure 7. The posting activity of the mailing list over time is also shown in Figure 7, averaging 164 posts per month since MoveIt! was launched.
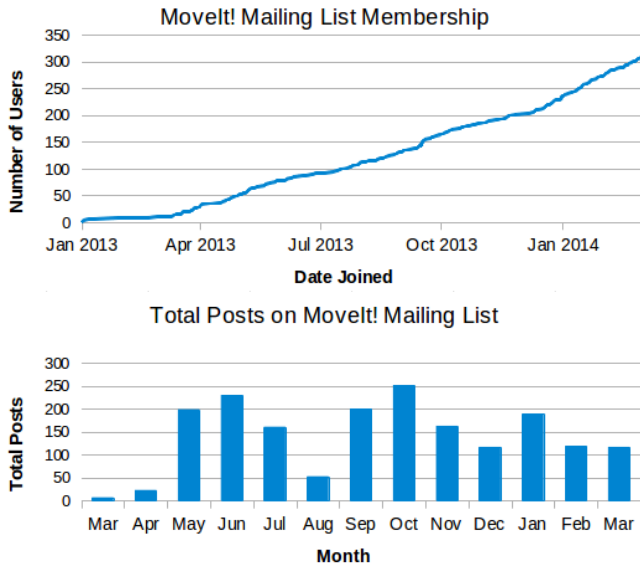
Fig. 7. MoveIt! mailing list statistics though March 2014. MoveIt! was alpha released in May 2013
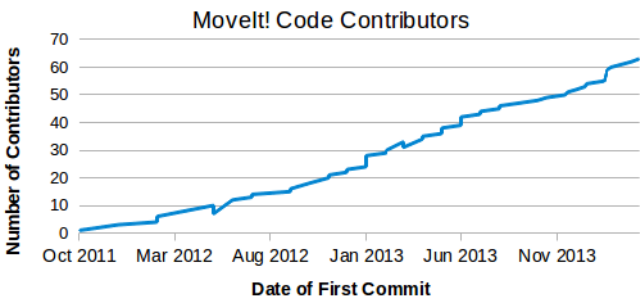

Fig. 8. MoveIt! source repository contributions through May 2014

There have been a total of 63 contributors to the MoveIt! code base since its initial development began in 2011. The total number of contributors over time is shown in Figure 8. According to statistics gathered by the website Ohloh.com, which tracks activity of open source projects, MoveIt! is "one the largest open-source teams in the world and is in the top 2% of all project teams on Ohloh" [39].

5.2 Comparison

A brief comparison with MoveIt! to OMPL and OpenRave is shown in Figure 9. In this diagram the total number of code contributors is plotted with respect to time, as reported from the projects respective version control system (VCS). No other software projects discussed in this paper had VCSs available publicly for comparison.
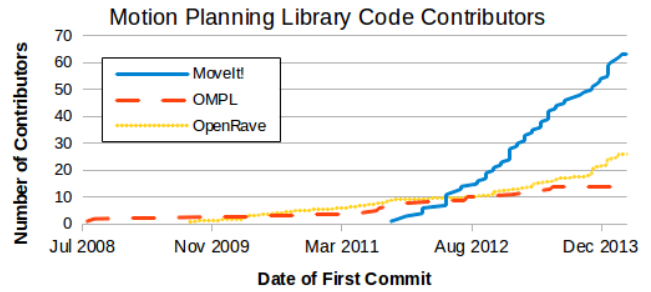

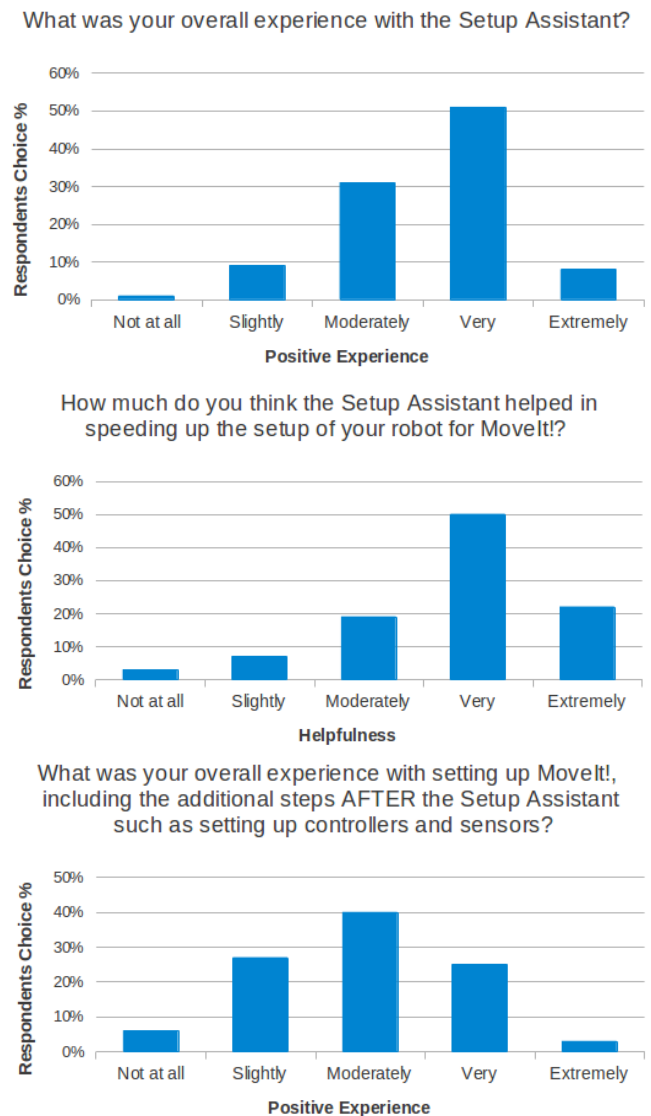Fig. 9. Comparison of code contributions between MoveIt!, OMPL, and OpenRave

Fig. 10. Survey data of 105 respondents on the MoveIt! and ROS mailing lists.

5.3 Survey

A survey on users' experience with MoveIt! was administered on the MoveIt! and ROS mailing lists. There were a total of 105 respondents; graduate students represented by far the largest group of respondents (39%), while faculty/post-docs (18%) and industry R&D users (17%) represented the next biggest groups (see [2] for the full survey results). Relevant results corresponding to the use of the MoveIt! Setup Assistant are shown in Figure 10.

Respondents were asked to rate their overall experience with using the MoveIt! SA, and asked how much the SA helped in speeding up setup of a robot in MoveIt!. For both questions, ninety percent had "moderately" to "extremely" positive experiences with the SA, and for both questions over half rated their experience as "very good."

Respondents then were asked what their overall experience was with setting up and configuring MoveIt!, including any additional steps they had to take after the SA, such as setting up controllers and sensors. In this question, the results were less positive. Forty percent had a "moderately" positive experience, but only 28% had a "very" or "extremely" positive experience.

An additional question asked respondents "how many minutes would you estimate you spent going from a URDF to solving motion plans using the MoveIt! Motion Planning Visualizer?" The responses to this question had a large amount of variance, with the mean time taking users 1.5 hours with a standard deviation of 2 hours.

6 DISCUSSION

6.1 MoveIt! Setup

We believe the barrier to entry for MoveIt! is easier than most, if not all, open source motion planning software available today as discussed in Section 2.1. As a result, MoveIt! has quickly become popular in the robotics community as a powerful MPF that is extensible to most users' needs for their robot application. The adoption rate of MoveIt! since its official release half a year ago has been very positive in comparison to the size of the world wide robotics community. With 309 users on the mailing list since MoveIt!'s release, a new member has joined the project at a rate of nearly one per day. These numbers indicate a healthy usage and popularity of this open source software project.

Community effort to improve MoveIt! has been better than expected given the large number of code contributors during MoveIt!'s existence. The comparison of two other robotics software project's all time contributors in Figure 9 makes it very evident that MoveIt! is a popular robotics project relative to others. Ohloh's ranking of MoveIt! as one of the largest open source teams in the world confirms our belief that by making complex software more accessible, more developers will be able to report and fix issues.

The results of the survey on MoveIt! indicated most people have found the cornerstone of our approach to lowering the barrier of entry, the Setup Assistant, to be very or extremely helpful in saving them time during setup (72% of respondents). Additionally, their overall experience was very or extremely positive with the SA (59%). However, in asking respondents their overall configuration experience with MoveIt! beyond just the SA, their ratings were lower, with only 28% saying they had a very or extremely positive experience setting up MoveIt!. This indicates that improvement can be made in the overall integration process and that adding more steps and features to the SA could reduce even further the entry barrier to MoveIt!.

From the lower results from this last survey question, it is clear that the setup and configuration process of MoveIt! can still be improved. A popular response from some of the free-form questions in the survey is that setting up the hardware controllers can also be a difficult task for non-experts, and the MoveIt! setup process does not yet document and provide example code as well as it could. It is likely this step will continue to require some custom coding to account for arbitrary hardware interfaces and communication methods, but based on the feedback we have received from actual users, this is certainly an area of improvement for the MoveIt! Setup Assistant to address.

The estimated setup time from taking a URDF and using MoveIt! to solve motion plans shows a large range of variance and likely indicates that wide range of experience levels in MoveIt! users. Although users averaged 1.5 hours to configuring MoveIt! from scratch, 31% reported it taking them 15 minutes or less to setup MoveIt! for a new robot. Creating software powerful but simple enough for all skill levels of users is a challenging task that MoveIt! will continue to tackle.

Though not exactly within the scope of MoveIt!, creating the robot model itself is a difficult task that typically requires a lot of trial and error in configuring the links and joints properly. This process could be improved by a better GUI for making arbitrary robot models, better tools for attaching the links together correctly, and more documentation.

Finally, although MoveIt! is very extensible with its plugin-based architecture, modifying the actual code base of MoveIt! can be intimidating due to its large size. MoveIt! contains over 170 thousand lines of code across all its various packages. Due to the need for computational speed and power, the layout of the code can sometimes seem complicated and abstracted.

We would like to emphasize the effect of a quick setup process and *Getting Started* demo on a new user unaccustomed to MoveIt! or motion planning in general. The positive reinforcement of a quick initial success encourages novices to continue to use the software and enables them to begin going deeper into the functionality and code base. If the entry barrier is too high, that is to say if it is too complex and error-prone, a new user will likely give up and turn to other frameworks or custom solutions. Attempting to blindly fix software that a

new user has not had any success with is a very difficult task.

6.2 MoveIt Development

Finding the balance between the opposing objectives of the *entry barrier design principles* was a difficult task in developing MoveIt!. *Immediacy* was given one of the highest priorities, such that we focused on users being able to go from robot model to planning feasible motion plans with very few steps. To allow this, the GUI streamlined the entire process and only presented the most important and *intuitive* configuration options. For example, the concept of defining the parts of a robot that make up an arm is very *intuitive*. This focus on *immediacy* sacrifices *transparency* in that once users get this initial virtual demonstration working, they have not learned much on how to extend or dive deeper into the MPF. At this point *documentation* is necessary to the user.

Another pair of conflicting principles are *extensibility* and *intuitiveness*. The powerful plugin framework that MoveIt! provides allows custom components to be loaded and swapped out at runtime. However, this requires many layers of abstraction and inheritance, and results in overall convoluted code that is difficult for new developers. The balance is attempted by providing documentation and code examples for plugins that allows users to build new components without worrying about the underlying framework.

Integrating components from different sources, such as third party libraries of robotic software from other research groups, presents challenges as discussed in [18]. During MoveIt! development, the plugin interfaces were required to be general enough to work with many different implementation methods and choices of data structures. This was accomplished by providing "wrapper" packages that connect together the standard MoveIt! plugin interfaces with the third party software API. For example, MoveIt is currently setup to work with at least three planning libraries – OMPL, SBPL [40], and CHOMP [41]. Although they represent fairly different approaches to motion planning and use different datastructures, each has a wrapper component that harmonizes them to work together in MoveIt!. It should be noted that as is true with any external dependency, maintaining compatibility with these wrappers has proven challenging.

6.3 Robotic Software

The techniques utilized in lowering the barrier of entry for MoveIt! can easily be applied to robotics software in general. Almost all robotics software requires customizations specific to a particular hardware and kinematic configuration. Reducing the difficulty of performing these customizations should be the goal of robotic software engineers who desire to create useful tools for a large audience.

For example, perception applications such as visual servoing require similar kinematic models to those being used in motion planning. Frame transforms must be specified for the location of the camera and the location of the end effector with respect to the rest of the robot's geometry [42]. This is often a difficult and tedious task. Automating the setup and calibration of these transforms lowers the barrier of entry to new users to robotic vision software and makes the vision software useful to more users. In general, automating the sequence of configuration steps necessary for performing particular tasks is a useful strategy: the users will not have to think about whether they have missed steps or whether they have performed the necessary steps in the correct order.

The *entry barrier design principles* of immediacy, transparency, intuitiveness, reconfigurability, extensibility, and documentation present a set of guidelines for other open source robotic software projects to reduce their barriers of entry to users. In fact, many existing robotics projects already follow subsets of these principles, but typically to a lesser extent and fervor.

Creating a GUI such as the Setup Assistant is a time consuming process that many robotics developers avoid in favor of hard-coded or command-line based configuration, thereby neglecting the opportunity to attract non-expert users. Between two developers, the various GUIs and configuration tools in MoveIt! took about three months of development time. We believe that the trade-off in the time invested is worthwhile for the higher adoption rates and creation of a larger community willing to contribute to the software's development.

It is understandable that when robotics research is the priority, spending time on tangential aspects of the project such as GUIs and configuration tools can be less important to the researcher. Still, we would like to encourage researchers and developers alike, when possible, to spend the extra time making their work reusable by taking into consideration the barriers to entry that other users might encounter. Too often, software is touted as "open source" when its usefulness is in actuality severely limited by the difficulties other users encountered in applying it to other robotic projects. By sharing accessible open source robotics software, the progress of robotic technology is accelerated and the robotics community as a whole benefits.

7 CONCLUSION

Beyond the usual considerations in building successful robotics software, an open source project that desires to maintain an active and large user base needs to take into account the barriers of entry to new users. By making robotic software more accessible, more users have the ability to utilize and contribute to robotics development who previously could not have. The entry barrier design principles are guidelines for robotic software engineers to improve the usefulness and usability of their work to others as demonstrated in this paper with the case study of MoveIt!.

As robotic algorithms become more complicated and the number of interacting software components and size of the

code base increases, configuring an arbitrary robot to utilize robotic software becomes a daunting task requiring domain-specific expertise in a very large breadth of theory and implementation. To account for this, quick and easy initial configuration, with partially automated optimization and easily extensible components for future customization, is becoming a greater necessity in motion planning and in robotic software engineering in general.

ACKNOWLEDGMENTS

The authors would like to acknowledge E. Gil Jones as the initiator and Matthew Klingensmith as the implementor of the Arm Navigation Setup Wizard, the inspiration for the MoveIt! Setup Assistant.

REFERENCES

[1] M. Quigley, K. Conley, B. Gerkey, J. Faust, T. Foote, J. Leibs, R. Wheeler, and A. Y. Ng, "ROS: an open-source robot operating system," in *ICRA Workshop On Open Source Software*, vol. 3, no. 3.2, 2009. 1

[2] I. A. Şucan and S. Chitta. (2013, Oct.) Moveit! [Online]. Available: http://moveit.ros.org/ 1, 1.5, 2.2, 4.1, 4.2, 4.3, 4.3, 5.1, 5.3

[3] H. Bruyninckx, "Open robot control software: the orocos project," in *Robotics and Automation, 2001. Proceedings 2001 ICRA. IEEE International Conference on*, vol. 3. IEEE, 2001, pp. 2523–2528. 1, 1.2, 1.4, 4.4

[4] A. Makarenko, A. Brooks, and T. Kaupp, "On the benefits of making robotic software frameworks thin," in *Intelligent Robots and Systems (IROS'07), 2007. IEEE International Conference on*. IEEE, 2007. 1, 1.2, 2.2

[5] J. M. Carroll, *Interfacing thought: Cognitive aspects of human-computer interaction*. The MIT Press, 1987. 1.1

[6] D. C. Schmidt, "Why software reuse has failed and how to make it work for you," *C++ Report*, vol. 11, no. 1, p. 1999, 1999. 1.2

[7] D. C. Schmidt and A. Porter, "Leveraging open-source communities to improve the quality & performance of open-source software," in *Proceedings of the 1st Workshop on Open Source Software Engineering*, 2001. 1.2

[8] N. Correll, R. Wing, and D. Coleman, "A one-year introductory robotics curriculum for computer science upperclassmen," *Education, IEEE Transactions on*, vol. 56, no. 1, pp. 54–60, 2013. 1.2

[9] M. Moll, I. A. Şucan, J. Bordeaux, and L. E. Kavraki, "Teaching motion planning concepts to undergraduate students," in *Advanced Robotics and its Social Impacts (ARSO), 2011 IEEE Workshop on*. IEEE, 2011, pp. 27–30. 1.2, 1.2

[10] L. Guyot, N. Heiniger, O. Michel, and F. Rohrer, "Teaching robotics with an open curriculum based on the e-puck robot, simulations and competitions," in *Proceedings of the 2nd International Conference on Robotics in Education (RiE 2011)*, 2011. 1.2

[11] K. J. Boudreau, "Let a thousand flowers bloom? an early look at large numbers of software app developers and patterns of innovation," *Organization Science*, vol. 23, no. 5, pp. 1409–1427, 2012. 1.2

[12] H. A. Yanco and J. L. Drury, "A taxonomy for human-robot interaction," in *Proceedings of the AAAI Fall Symposium on Human-Robot Interaction*, 2002, pp. 111–119. 1.2, 1.4, 4.1

[13] W. D. Smart, "Is a common middleware for robotics possible?" in *Proceedings of the IROS 2007 workshop on Measures and Procedures for the Evaluation of Robot Architectures and Middleware*. Citeseer, 2007. 1.3

[14] S. Kchir, T. Ziadi, M. Ziane, and S. Stinckwich, "A top-down approach to managing variability in robotics algorithms," *arXiv preprint arXiv:1312.7572*, 2013. 1.3, 1.4

[15] Y.-H. Kuo and B. MacDonald, "A distributed real-time software framework for robotic applications," in *Robotics and Automation, 2005. ICRA 2005. Proceedings of the 2005 IEEE International Conference on*, 2005, pp. 1964–1969. 1.4

[16] T. H. Collett, B. A. MacDonald, and B. P. Gerkey, "Player 2.0: Toward a practical robot programming framework," in *Australasian Conference on Robotics and Automation*, 2005. 1.4

[17] J. Kramer and M. Scheutz, "Development environments for autonomous mobile robots: A survey," *Autonomous Robots*, vol. 22, no. 2, pp. 101–132, 2007. 1.4

[18] D. Brugali, W. Nowak, L. Gherardi, A. Zakharov, and E. Prassler, "Component-based refactoring of motion planning libraries," in *Intelligent Robots and Systems (IROS), 2010 IEEE/RSJ International Conference on*. IEEE, 2010, pp. 4042–4049. 1.4, 6.2

[19] A. Steinfeld, T. Fong, D. Kaber, M. Lewis, J. Scholtz, A. Schultz, and M. Goodrich, "Common metrics for human-robot interaction," in *Proceedings of the 1st ACM SIGCHI/SIGART conference on Human-robot interaction*, ser. HRI '06. New York, NY, USA: ACM, 2006, pp. 33–40. [Online]. Available: http://doi.acm.org/10.1145/1121241.1121249 1.4

[20] M. A. Goodrich and D. R. Olsen Jr, "Seven principles of efficient human robot interaction," in *Systems, Man and Cybernetics, 2003. IEEE International Conference on*, vol. 4. IEEE, 2003, pp. 3942–3948. 1.4

[21] M. W. Kadous, R. K.-M. Sheh, and C. Sammut, "Effective user interface design for rescue robotics," in *Proceedings of the 1st ACM SIGCHI/SIGART conference on Human-robot interaction*, ser. HRI '06. New York, NY, USA: ACM, 2006, pp. 250–257. [Online]. Available: http://doi.acm.org/10.1145/1121241.1121285 1.4

[22] S. Chitta, E. G. Jones, M. Ciocarlie, and K. Hsiao, "Perception, planning, and execution for mobile manipulation in unstructured environments," *IEEE Robotics and Automation Magazine*, vol. 19, no. 2, pp. 58–71, 2012. 1.4, 2.1

[23] A. Perez and J. Rosell, "A roadmap to robot motion planning software development," *Computer Applications in Engineering Education*, vol. 18, no. 4, pp. 651–660, 2010. 2

[24] S. LaValle, P. Cheng, J. Kuffner, S. Lindemann, A. Manohar, B. Tovar, L. Yang, and A. Yershova. (2013, Oct.) Msl: Motion strategy library. [Online]. Available: http://msl.cs.uiuc.edu/msl/ 2.1

[25] F. Schwarzer, M. Saha, and J. Latombe. (2013, Nov.) Software: Motion planning kit. [Online]. Available: http://robotics.stanford.edu/~mitul/mpk/ 2.1

[26] E. Plaku, K. Bekris, and E. Kavraki, "Oops for motion planning: An online, open-source, programming system," in *Robotics and Automation, 2007 IEEE International Conference on*, 2007, pp. 3711–3716. 2.1

[27] I. A. Şucan, M. Moll, and L. E. Kavraki, "The Open Motion Planning Library," *IEEE Robotics & Automation Magazine*, vol. 19, no. 4, pp. 72–82, December 2012, http://ompl.kavrakilab.org. 2.1

[28] R. Diankov and J. Kuffner, "Openrave: A planning architecture for autonomous robotics," *Robotics Institute, Pittsburgh, PA, Tech. Rep. CMU-RI-TR-08-34*, p. 79, 2008. 2.1

[29] (2013, Oct.) Iso/pas 17506:2012 industrial automation systems and integration – collada digital asset schema specification for 3d visualization of industrial data. [Online]. Available: http://www.iso.org/iso/catalogue_detail.htm?csnumber=59902 2.1, 4.1

[30] K. A. Wyrobek, E. H. Berger, H. M. Van der Loos, and J. K. Salisbury, "Towards a personal robotics development platform: Rationale and design of an intrinsically safe personal robot," in *Robotics and Automation, 2008. ICRA 2008. IEEE International Conference on*. IEEE, 2008, pp. 2165–2170. 2.2

[31] R. O. Ambrose, H. Aldridge, R. S. Askew, R. R. Burridge, W. Bluethmann, M. Diftler, C. Lovchik, D. Magruder, and F. Rehnmark, "Robonaut: Nasa's space humanoid," *Intelligent Systems and their Applications, IEEE*, vol. 15, no. 4, pp. 57–63, 2000. 2.2

[32] J. Pan, S. Chitta, and D. Manocha, "Fcl: A general purpose library for collision and proximity queries," in *Robotics and Automation (ICRA), 2012 IEEE International Conference on*, 2012, pp. 3859–3866. 2.2

[33] R. Smits. (2013, Oct.) Kdl: Kinematics and dynamics library. [Online]. Available: http://www.orocos.org/kdl 2.2

[34] W. O. Galitz, *The essential guide to user interface design: an introduction to GUI design principles and techniques*. Wiley, 2007. 3

[35] A. Forward and T. C. Lethbridge, "The relevance of software documentation, tools and technologies: a survey," in *Proceedings of the 2002 ACM symposium on Document engineering*. ACM, 2002, pp. 26–33. 3

[36] W. Garage. (2013, Oct.) Urdf: Universal robotic description format. [Online]. Available: http://wiki.ros.org/urdf 4.1

[37] B. Cohen, I. A. Şucan, and S. Chitta, "A generic infrastructure for benchmarking motion planners," in *Intelligent Robots and Systems (IROS), 2012 IEEE/RSJ International Conference on*. IEEE, 2012, pp. 589–595. 4.2

[38] R. Diankov. (2013, Oct.) Ikfast: The robot kinematics compiler. [Online]. Available: http://openrave.org/docs/latest_stable/openravepy/ikfast/#ikfast-the-robot-kinematics-compiler 4.3

[39] Ohloh. (2013, Oct.) Ohloh: Moveit! project summary factoids. [Online]. Available: https://www.ohloh.net/p/moveit_/factoids#FactoidTeamSizeVeryLarge 5.1

[40] M. Likhachev. Sbpl graph search library. 6.2

[41] N. Ratliff, M. Zucker, J. A. Bagnell, and S. Srinivasa, "Chomp: Gradient optimization techniques for efficient motion planning," in *Robotics and Automation, 2009. ICRA'09. IEEE International Conference on*. IEEE, 2009, pp. 489–494. 6.2

[42] S. Hutchinson, G. Hager, and P. Corke, "A tutorial on visual servo control," *Robotics and Automation, IEEE Transactions on*, vol. 12, no. 5, pp. 651–670, 1996. 6.3

Nikolaus Correll , M.S., 2003, Ph.D., 2007, is an Assistant Professor in Computer Science at the University of Colorado. He obtained his Master's degree in Electrical Engineering from ETH Zürich, a Ph.D. in Computer Science from EPFL and worked as a post-doc at MIT CSAIL. Nikolaus's research interests are multi-robot and swarming systems. He is the recipient of a 2012 NSF CAREER and a 2012 NASA Early Career Faculty fellowship.

David Coleman , M.S., 2013, is a Ph.D. candidate in Computer Science at the University of Colorado. He obtained his B.S. in Mechanical Engineering at Georgia Tech and Masters in C.S. at CU Boulder. David's research interests include robotic manipulation, motion planning, and controls. He is one of the main developers of MoveIt!, creating the MoveIt! Setup Assistant while interning at Willow Garage continuing development at the Open Source Robotics Foundation and with his regular research.

Ioan A. Şucan , received the Ph.D. degree in computer science at Rice University, Houston, TX, in 2011. His research interests include mobile manipulation, task and motion planning for manipulation and motion planning under differential constraints. While at Willow Garage Ioan initiated the core components that make up the Arm Navigation software platform, started the MoveIt! software platform and the Open Motion Planning Library (OMPL). Ioan is the main developer of MoveIt! and OMPL.

Sachin Chitta , Ph.D., is associate director of robotics systems and software in the Robotics Program at SRI International. His research interests include mobile manipulation, motion planning and learning for manipulation. Sachin Chitta was at Willow Garage from 2007-2013 and was a core member of the team that developed the PR2 robot and the Robot Operating System (ROS). He initiated and led the development of the MoveIt! and Arm Navigation software platforms to enable advanced manipulation capabilities for any robot. He obtained his PhD from the Grasp Lab at the University of Pennsylvania in 2005.